

Capstone Project - TESLA's Energy Performance Analysis Model

October 29, 2023

```
[1]: #Importing All major The Necessary Packages
import numpy as np
import pandas as pd
from pandas import DataFrame, Series
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MaxAbsScaler
```

```
[21]: dff =pd.read_csv(r'C:\Users\Amardeep\Downloads\1.electric_vehicle_data.csv')
```

```
[22]: dff
```

```
[22]:
```

	Vehicle_ID	Battery_Capacity_kWh	Motor_Power_kW	Weight_kg	\
0	1	91.453287	144.713077	2306.402444	
1	2	55.596029	278.426210	2433.488911	
2	3	49.334685	141.111622	1207.229380	
3	4	43.977956	252.662106	1595.065551	
4	5	94.117021	100.060068	1279.696413	
..	
995	996	83.403480	194.603532	2030.196131	
996	997	68.448818	198.239941	1287.598985	
997	998	69.803940	246.440589	1482.724663	
998	999	58.350458	60.074563	1588.490647	
999	1000	94.693425	81.283579	2007.158867	

	Temperature_Celsius	Terrain_Type	Distance_Traveled_km	\
0	-9.817508	Highway	262.523326	
1	32.894289	Suburban	66.140868	
2	22.471884	Suburban	281.077227	
3	-4.788043	Highway	177.042720	
4	38.362606	Suburban	119.862466	
..	
995	1.780901	Urban	141.962617	
996	16.688313	Suburban	183.645218	
997	29.096608	Urban	134.730399	

998	33.446197	Urban	207.295620
999	17.217717	Highway	193.550225

	Average_Speed_kmph	Charging_Type	Energy_Consumption_kWh_per_100km
0	16.858933	Fast	806.134311
1	38.096844	Home	1214.442507
2	38.017746	Home	240.234773
3	55.472865	Home	315.630163
4	33.800371	Home	1661.105114
..
995	98.436035	Home	1259.114483
996	44.428797	Home	884.890093
997	39.790435	Fast	1217.302617
998	69.097834	Home	536.247522
999	22.226434	Home	679.260536

[1000 rows x 10 columns]

```
[23]: one_hot=pd.get_dummies(dff['Terrain_Type'])
      # Drop column as it is now encoded
      dff = dff.drop('Terrain_Type',axis = 1)
      # Join the encoded df
      dff = dff.join(one_hot)
```

```
[25]: two_hot=pd.get_dummies(dff['Charging_Type'])
      # Drop column as it is now encoded
      dff = dff.drop('Charging_Type',axis = 1)
      # Join the encoded df
      dff = dff.join(two_hot)
```

```
[26]: dff
```

```
[26]:
```

	Vehicle_ID	Battery_Capacity_kWh	Motor_Power_kW	Weight_kg	\
0	1	91.453287	144.713077	2306.402444	
1	2	55.596029	278.426210	2433.488911	
2	3	49.334685	141.111622	1207.229380	
3	4	43.977956	252.662106	1595.065551	
4	5	94.117021	100.060068	1279.696413	
..	
995	996	83.403480	194.603532	2030.196131	
996	997	68.448818	198.239941	1287.598985	
997	998	69.803940	246.440589	1482.724663	
998	999	58.350458	60.074563	1588.490647	
999	1000	94.693425	81.283579	2007.158867	

	Temperature_Celsius	Distance_Traveled_km	Average_Speed_kmph	\
0	-9.817508	262.523326	16.858933	

1	32.894289	66.140868	38.096844
2	22.471884	281.077227	38.017746
3	-4.788043	177.042720	55.472865
4	38.362606	119.862466	33.800371
..
995	1.780901	141.962617	98.436035
996	16.688313	183.645218	44.428797
997	29.096608	134.730399	39.790435
998	33.446197	207.295620	69.097834
999	17.217717	193.550225	22.226434

	Energy_Consumption_kWh_per_100km	Highway	Suburban	Urban	Fast	Home	\
0	806.134311	1	0	0	1	0	
1	1214.442507	0	1	0	0	1	
2	240.234773	0	1	0	0	1	
3	315.630163	1	0	0	0	1	
4	1661.105114	0	1	0	0	1	
..	
995	1259.114483	0	0	1	0	1	
996	884.890093	0	1	0	0	1	
997	1217.302617	0	0	1	1	0	
998	536.247522	0	0	1	0	1	
999	679.260536	1	0	0	0	1	

	Public
0	0
1	0
2	0
3	0
4	0
..	...
995	0
996	0
997	0
998	0
999	0

[1000 rows x 14 columns]

```
[27]: E = dff
```

```
[28]: E.index = dff['Vehicle_ID']
E = E.drop(columns=['Vehicle_ID'])
```

```
[29]: #Scaling the bin and obtaining the scores for the pca analysis.
E_index=E.index
E= StandardScaler().fit_transform(E)
```

```
pca = PCA(n_components=3, svd_solver='full')
PC_scores = pca.fit_transform(E)
scores_pd = pd.DataFrame(data = PC_scores
                          ,columns = ['PC1', 'PC2', 'PC3']
                          ,index =E_index)
```

[32]: A2

```
[32]:
```

	Battery_Capacity_kWh	Motor_Power_kW	Weight_kg \
Vehicle_ID			
1	91.453287	144.713077	2306.402444
2	55.596029	278.426210	2433.488911
3	49.334685	141.111622	1207.229380
4	43.977956	252.662106	1595.065551
5	94.117021	100.060068	1279.696413
...
996	83.403480	194.603532	2030.196131
997	68.448818	198.239941	1287.598985
998	69.803940	246.440589	1482.724663
999	58.350458	60.074563	1588.490647
1000	94.693425	81.283579	2007.158867

	Temperature_Celsius	Distance_Traveled_km	Average_Speed_kmph \
Vehicle_ID			
1	-9.817508	262.523326	16.858933
2	32.894289	66.140868	38.096844
3	22.471884	281.077227	38.017746
4	-4.788043	177.042720	55.472865
5	38.362606	119.862466	33.800371
...
996	1.780901	141.962617	98.436035
997	16.688313	183.645218	44.428797
998	29.096608	134.730399	39.790435
999	33.446197	207.295620	69.097834
1000	17.217717	193.550225	22.226434

	Highway	Suburban	Urban	Fast	Home	Public
Vehicle_ID						
1	1	0	0	1	0	0
2	0	1	0	0	1	0
3	0	1	0	0	1	0
4	1	0	0	0	1	0
5	0	1	0	0	1	0
...
996	0	0	1	0	1	0
997	0	1	0	0	1	0
998	0	0	1	1	0	0

999	0	0	1	0	1	0
1000	1	0	0	0	1	0

[1000 rows x 12 columns]

[33]: *#Dropping unnecessary columns from the dataset.*

```
A2=dff.drop(columns=['Vehicle_ID'])
loadings_pd = pd.DataFrame(data = pca.components_.T
                           ,columns = ['PC1', 'PC2', 'PC3']
                           ,index = A2.columns)
```

[34]: *#Plotting the PCA plot function*

```
def myplot(scores,loadings,loading_labels=None,score_labels=None):
    # adjusting the scores to fit in (-1,1)
    xt = scores[:,0]
    yt = scores[:,1]
    n = loadings.shape[0]
    scalext = 1.0/(xt.max() - xt.min())
    scaleyt = 1.0/(yt.max() - yt.min())
    xt_scaled = xt * scalext
    yt_scaled = yt * scaleyt
    # adjusting the loadings to fit in (-1,1)
    p = loadings
    p_scaled = MaxAbsScaler().fit_transform(p)


    #plt.scatter(xs * scalex,ys * scaley, s=10)
    for (x,y), label in zip(np.vstack((xt_scaled, yt_scaled)).T,score_labels):
        plt.text(x, y, label, ha='center', size=11)

    for i in range(n):
        plt.arrow(0, 0, p_scaled[i,0], p_scaled[i,1], color = 'r',alpha = 0.5)
        if loading_labels is None:
            plt.text(p_scaled[i,0], p_scaled[i,1], "Var"+str(i+1), color = 'g',
            ↪ha = 'center', va = 'center')
        else:
            plt.text(p_scaled[i,0], p_scaled[i,1], loading_labels[i], color =
            ↪'r', ha = 'center', va = 'center', size=16)
    plt.xlim(-1,1)
    plt.ylim(-1,1)
    plt.xlabel("PC{}".format(1), fontsize=20);
    plt.ylabel("PC{}".format(2), fontsize=20);
    plt.tick_params(labelsize=16)
    plt.grid()
```

[35]: plt.rcParams["figure.figsize"] = [16,9]

```
myplot(PC_scores[:, :2],loadings_pd.iloc[:, :2],loading_labels=loadings_pd.
↪index,score_labels=scores_pd.index)
```

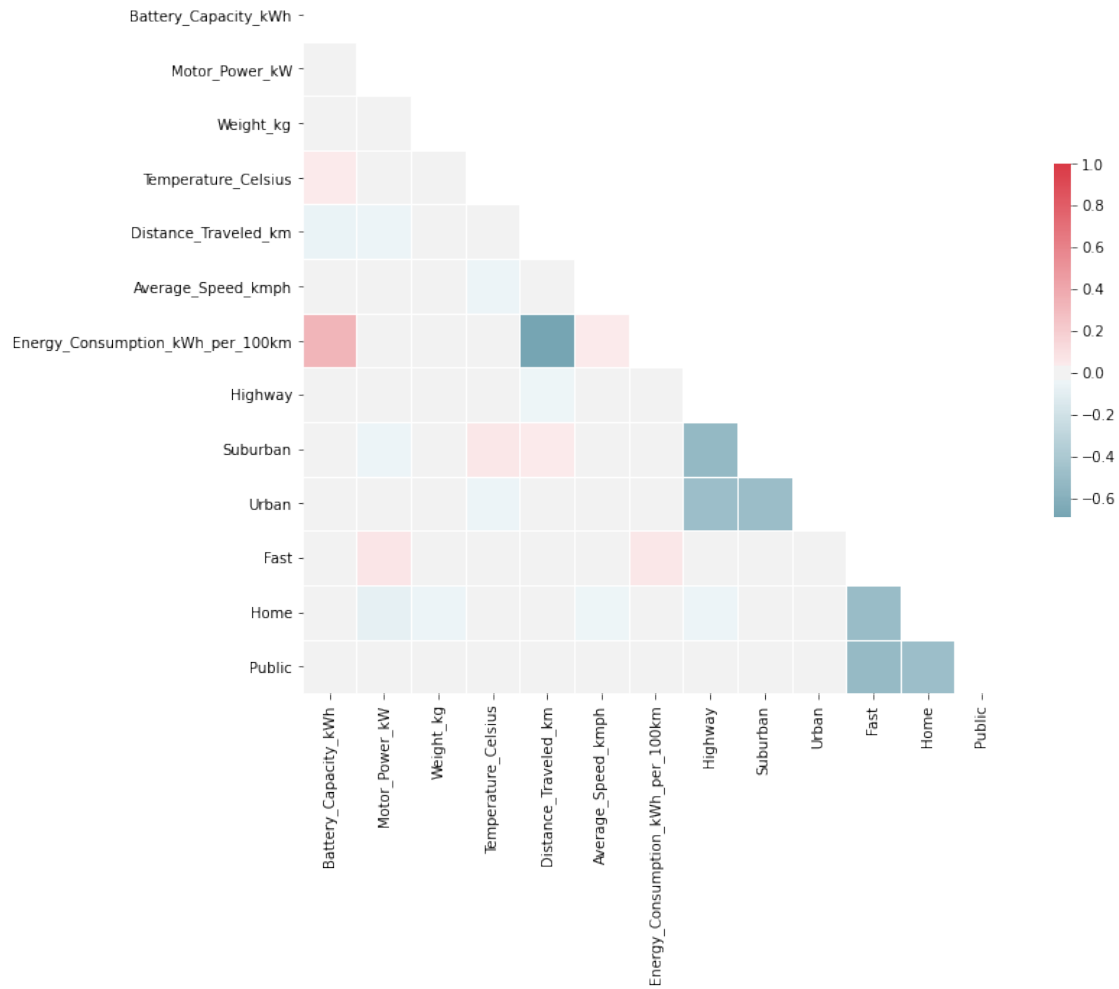
```
plt.show()
```



```
[36]: #Creating a correlation plot for the variables
import seaborn as sns
corr = A2.corr()
# Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=1, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
plt.show()
plt.figure(figsize=(50,50))
```




[36]: <Figure size 3600x3600 with 0 Axes>

<Figure size 3600x3600 with 0 Axes>


[106]: A2.hist(column='Battery_Capacity_kWh')

[106]: array([[<AxesSubplot:title={'center': 'Battery_Capacity_kWh'}>]],
dtype=object)




```
[107]: #Scatter plots between pm2.5 and pressure.
plt.scatter(x=A2['Motor_Power_kW'],y=A2['Energy_Consumption_kWh_per_100km'])
```

```
[107]: <matplotlib.collections.PathCollection at 0x23194f0a708>
```




```
[ ]:
```

```
[37]: sns.pairplot(A2)
plt.show()
```




```
[38]: #Performing PCA for 10 components
pca10 = PCA(n_components=10, svd_solver='full')
pca10.fit(E)
```

```
[38]: PCA(n_components=10, svd_solver='full')
```

```
[39]: #Variance for components displayed in graph.
plt.plot(range(1,pca10.n_components+1), np.cumsum(pca10.
    →explained_variance_ratio_),'-o')
plt.xlabel('Number of components', fontsize=20)
plt.ylabel('Cumulative explained variance', fontsize=20);
plt.tick_params(labelsize=16)
plt.ylim(0,1)

plt.show()
```



```
[41]: A2
```

```
[41]:
```

Vehicle_ID	Battery_Capacity_kWh	Motor_Power_kW	Weight_kg \
1	91.453287	144.713077	2306.402444
2	55.596029	278.426210	2433.488911
3	49.334685	141.111622	1207.229380
4	43.977956	252.662106	1595.065551
5	94.117021	100.060068	1279.696413
...
996	83.403480	194.603532	2030.196131
997	68.448818	198.239941	1287.598985
998	69.803940	246.440589	1482.724663
999	58.350458	60.074563	1588.490647
1000	94.693425	81.283579	2007.158867

Vehicle_ID	Temperature_Celsius	Distance_Traveled_km	Average_Speed_kmph	\
1	-9.817508	262.523326	16.858933	
2	32.894289	66.140868	38.096844	
3	22.471884	281.077227	38.017746	
4	-4.788043	177.042720	55.472865	
5	38.362606	119.862466	33.800371	
...	
996	1.780901	141.962617	98.436035	
997	16.688313	183.645218	44.428797	
998	29.096608	134.730399	39.790435	
999	33.446197	207.295620	69.097834	
1000	17.217717	193.550225	22.226434	

Vehicle_ID	Energy_Consumption_kWh_per_100km	Highway	Suburban	Urban	Fast	\
1	806.134311	1	0	0	1	
2	1214.442507	0	1	0	0	
3	240.234773	0	1	0	0	
4	315.630163	1	0	0	0	
5	1661.105114	0	1	0	0	
...	
996	1259.114483	0	0	1	0	
997	884.890093	0	1	0	0	
998	1217.302617	0	0	1	1	
999	536.247522	0	0	1	0	
1000	679.260536	1	0	0	0	

Vehicle_ID	Home	Public
1	0	0
2	1	0
3	1	0
4	1	0
5	1	0
...
996	1	0
997	1	0
998	0	0
999	1	0
1000	1	0

[1000 rows x 13 columns]

```
[42]: #Splitting the input and output into separate dataframes
y=A2['Energy_Consumption_kWh_per_100km']
X=A2.drop(columns=['Energy_Consumption_kWh_per_100km'])
```

```
[47]: #Standard test using normal conditions
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.60)
#Create a Gaussian Classifier
rf1=RandomForestRegressor(n_estimators = 100, max_depth=2)

#Train the model using the training sets y_pred=clf.predict(X_test)
rf1.fit(X_train,y_train)

# prediction on test set
y_pred=rf1.predict(X_test)
```

```
[50]: from sklearn.linear_model import LarsCV
from sklearn.datasets import make_regression
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.55,
    ↪random_state=40)
reg = LarsCV(cv=10).fit(X_train, y_train)
reg.score(X_train, y_train)
```

```
[50]: 0.589415510989831
```

```
[51]: reg.score(X_test,y_test)
```

```
[51]: 0.5382427724098321
```

```
[55]: def get_params(self, deep=True):
    """Get parameters for this estimator.
    Parameters
    -----
    deep : boolean, optional
        If True, will return the parameters for this estimator and
        contained subobjects that are estimators.
    Returns
    -----
    params : mapping of string to any
        Parameter names mapped to their values.
    """
    out = dict()
    for key in self._get_param_names():
        value = getattr(self, key, None)
        if deep and hasattr(value, 'get_params'):
            deep_items = value.get_params().items()
            out.update((key + '__' + k, val) for k, val in deep_items)
        out[key] = value
    return out
```

```

def set_params(self, **params):
    """Set the parameters of this estimator.
    The method works on simple estimators as well as on nested objects
    (such as pipelines). The latter have parameters of the form
    ``<component>__<parameter>`` so that it's possible to update each
    component of a nested object.
    Returns
    -----
    self
    """
    if not params:
        # Simple optimization to gain speed (inspect is slow)
        return self
    valid_params = self.get_params(deep=True)

    nested_params = defaultdict(dict) # grouped by prefix
    for key, value in params.items():
        key, delim, sub_key = key.partition('__')
        if key not in valid_params:
            raise ValueError('Invalid parameter %s for estimator %s. '
                              'Check the list of available parameters '
                              'with `estimator.get_params().keys()`.`' %
                              (key, self))

        if delim:
            nested_params[key][sub_key] = value
        else:
            setattr(self, key, value)
            valid_params[key] = value

    for key, sub_params in nested_params.items():
        valid_params[key].set_params(**sub_params)

    return self

```

```

[57]: from sklearn import linear_model
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.70,
↳ random_state=40)
reg5 = linear_model.Ridge(alpha=.5)
reg5.fit(X,y)
reg5.coef_

```

```

[57]: array([ 17.35019792, -0.18979142, -0.08137701, -1.49856252,
-9.83798441,  1.03330851, -37.56926936,  32.08284499,
 5.48642436,  54.96204272, -17.01617768, -37.94586504])

```

```
[58]: from sklearn import linear_model
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.60,
↳random_state=40)
regbr1 = linear_model.BayesianRidge()
regbr1.fit(X_train, y_train)
regbr1.score(X_train, y_train)
```

[58]: 0.5871416296606013

```
[59]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.70,
↳random_state=40)
regbr2 = linear_model.BayesianRidge()
regbr1.fit(X_test, y_test)
regbr1.score(X_test, y_test)
```

[59]: 0.5658037238185146

```
[60]: from sklearn.linear_model import ElasticNetCV
from sklearn.datasets import make_regression
X, y = make_regression(n_features=10, random_state=0)
regr15 = ElasticNetCV(cv=5, random_state=0)
regr15.fit(X_train, y_train)
```

[60]: ElasticNetCV(cv=5, random_state=0)

```
[61]: regr15.fit(X_test, y_test)
regr15.score(X_test, y_test)
```

[61]: 0.5636727815316841

```
[62]: from sklearn.linear_model import OrthogonalMatchingPursuit
from sklearn.datasets import make_regression
X, y = make_regression(noise=4, random_state=0)
rego = OrthogonalMatchingPursuit().fit(X_train, y_train)
rego.score(X_train, y_train)
```

[62]: 0.4738243100664945

```
[63]: rego.fit(X_test, y_test)
rego.score(X_test, y_test)
```

[63]: 0.4768603067156493

```
[65]: #Code for demographic dataset II(Predicting Mean Debt)
import xgboost as xgb
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
```

```
data_dmatrix = xgb.DMatrix(data=X,label=y)
```

```
[67]: xg_reg = xgb.XGBRegressor(objective='reg:linear', colsample_bytree = 0.3, ↵  
      ↪learning_rate = 0.5,  
      max_depth = 10, alpha = 10, n_estimators = 15)
```

```
[68]: xg_reg.fit(X_train,y_train)  
      preds = xg_reg.predict(X_test)
```

[09:13:43] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.1.0/src/objective/regression_obj.cu:170: reg:linear is now deprecated in favor of reg:squarederror.

[09:13:43] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.1.0/src/objective/regression_obj.cu:170: reg:linear is now deprecated in favor of reg:squarederror.

```
[69]: rmse = np.sqrt(mean_squared_error(y_test, preds))  
      print("RMSE: %f" % (rmse))
```

RMSE: 116.070966

```
[70]: # evaluate adaboost ensemble for regression  
      from numpy import mean  
      from numpy import std  
      from sklearn.datasets import make_regression  
      from sklearn.model_selection import cross_val_score  
      from sklearn.model_selection import RepeatedKFold  
      from sklearn.ensemble import AdaBoostRegressor  
      from sklearn.model_selection import train_test_split  
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)  
      # define the model  
      model = AdaBoostRegressor()  
      # evaluate the model  
      cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)  
      n_scores = cross_val_score(model, X, y, scoring='neg_mean_absolute_error', ↵  
      ↪cv=cv, n_jobs=-1, error_score='raise')  
      # report performance  
      print('MAE: %.3f (%.3f)' % (mean(n_scores), std(n_scores)))
```

MAE: -85.363 (16.293)

```
[71]: stat_scaled = StandardScaler().fit_transform(E)
```

```
[72]: stat_scaled
```

```
[72]: array([[ 1.30295661, -0.37527277,  1.31224387, ...,  1.35978044,  
            -0.6828438 , -0.70339769],  
          [-0.44735753,  1.49614042,  1.60638955, ..., -0.73541284,  
            1.46446375, -0.70339769],
```

```

[-0.75299492, -0.42567777, -1.23182725, ..., -0.73541284,
 1.46446375, -0.70339769],
...,
[ 0.24617872,  1.04847823, -0.59418463, ...,  1.35978044,
 -0.6828438 , -0.70339769],
[-0.3129045 , -1.55985074, -0.34938589, ..., -0.73541284,
 1.46446375, -0.70339769],
[ 1.46111874, -1.26301504,  0.61963506, ..., -0.73541284,
 1.46446375, -0.70339769]])

```

```
[76]: sc=pd.DataFrame(data=stat_scaled,columns=A2.columns)
```

```

[77]: # evaluate adaboost ensemble for regression
y = sc['Energy_Consumption_kWh_per_100km']
X = sc.drop(columns=['Energy_Consumption_kWh_per_100km'])
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
# define the model
model = AdaBoostRegressor()
# evaluate the model
cv = RepeatedKfold(n_splits=10, n_repeats=3, random_state=1)
n_scores = cross_val_score(model, X, y, scoring='neg_mean_absolute_error',
    ↳cv=cv, n_jobs=-1, error_score='raise')
# report performance
print('MAE: %.3f (%.3f)' % (mean(n_scores), std(n_scores)))

```

MAE: -0.282 (0.038)

```

[79]: from sklearn.ensemble import AdaBoostRegressor
from sklearn import tree
from sklearn.model_selection import GridSearchCV
for depth in range (1,10):
    tree_regressor=tree.DecisionTreeRegressor(max_depth=depth,random_state=1)
    if tree_regressor.fit(X,y).tree_.max_depth<depth:
        break
    score=np.
    ↳mean(cross_val_score(tree_regressor,X,y,scoring='neg_mean_squared_error',
    ↳cv=10,n_jobs=1))
    print(depth, score)

```

```

1 -0.4665006452342194
2 -0.3298919813850067
3 -0.24338425905946942
4 -0.1844261002642531
5 -0.18390516798846918
6 -0.21473861072072667
7 -0.22093608635146839
8 -0.22545316537757892

```


9 -0.2480393922465834

```
[82]: from sklearn.tree import DecisionTreeRegressor
      from sklearn.ensemble import AdaBoostRegressor
```

```
# Fit regression model
```

```
regr_1 = DecisionTreeRegressor(max_depth=4)
```

```
regr_2 = AdaBoostRegressor(DecisionTreeRegressor(max_depth=4),
                           n_estimators=300)
```

```
regr_1.fit(X, y)
```

```
regr_2.fit(X, y)
```

```
# Predict
```

```
y_1 = regr_1.predict(X)
```

```
y_2 = regr_2.predict(X)
```

```
[86]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=0)
```

```
[87]: from sklearn.ensemble import RandomForestRegressor
```

```
regressor = RandomForestRegressor(n_estimators=20, random_state=0)
```

```
regressor.fit(X_train, y_train)
```

```
y_pred = regressor.predict(X_test)
```

```
[88]: from sklearn import metrics
```

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
```

```
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
```

```
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
```

Mean Absolute Error: 0.23470168723867146

Mean Squared Error: 0.19159806299984872

Root Mean Squared Error: 0.4377191599642957

```
[89]: from sklearn.linear_model import LarsCV
      from sklearn.datasets import make_regression
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.55,
      ↪random_state=40)
```

```
reg = LarsCV(cv=10).fit(X_train, y_train)
```

```
reg.score(X_train, y_train)
```

```
[89]: 0.5894155109898311
```

```
[90]: #Code for demographic dataset II(Predicting Mean Debt)
import xgboost as xgb
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
data_dmatrix = xgb.DMatrix(data=X,label=y)

[91]: xg_reg = xgb.XGBRegressor(objective='reg:linear', colsample_bytree = 0.3,
    ↪learning_rate = 0.5,
        max_depth = 10, alpha = 10, n_estimators = 15)

[92]: xg_reg.fit(X_train,y_train)
preds = xg_reg.predict(X_test)
```

[09:34:46] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.1.0/src/objective/regression_obj.cu:170: reg:linear is now deprecated in favor of reg:squarederror.
[09:34:46] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.1.0/src/objective/regression_obj.cu:170: reg:linear is now deprecated in favor of reg:squarederror.

```
[93]: rmse = np.sqrt(mean_squared_error(y_test, preds))
print("RMSE: %f" % (rmse))
```

RMSE: 0.522921

```
[98]: from math import sqrt
from sklearn.metrics import mean_squared_error,accuracy_score,r2_score
from sklearn.svm import SVR

# creating the model
svrmodel = SVR()

# feeding the training data to the model
svrmodel.fit(X_train, y_train)

# predicting the test set results
y_pred = svrmodel.predict(X_test)

# calculating the mean squared error
MSE = mean_squared_error(y_test, y_pred)
print('MSE: ', MSE)

# Calculating the root mean squared error
RMSE = sqrt(mean_squared_error(y_test, y_pred))
print('RMSE: ', RMSE)

# Calculating the r2 score
```

```
R2 = r2_score(y_test, y_pred)
print('R2 :', R2)
```

```
MSE: 0.3923297581109713
RMSE: 0.6263623217523315
R2 : 0.6306796061222809
```

```
[97]: from sklearn.tree import DecisionTreeRegressor
      # creating the model
      dtrmodel = DecisionTreeRegressor()

      # feeding the training data to the model
      dtrmodel.fit(X_train, y_train)

      # predicting the test set results
      y_pred = dtrmodel.predict(X_test)

      # calculating the mean squared error
      MSE = mean_squared_error(y_test, y_pred)
      print('MSE: ', MSE)

      # Calculating the root mean squared error
      RMSE = sqrt(mean_squared_error(y_test, y_pred))
      print('RMSE: ', RMSE)

      # Calculating the r2 score
      R2 = r2_score(y_test, y_pred)
      print('R2 :', R2)
```

```
MSE: 0.29684749698050544
RMSE: 0.5448371288564183
R2 : 0.7205620215139386
```

```
[99]: # creating the model
      #rfrmodel = RandomForestRegressor()
      rfrmodel = RandomForestRegressor(n_estimators = 40, max_depth = 4, n_jobs = -1)

      # feeding the training data to the model
      rfrmodel.fit(X_train, y_train)

      # predicting the test set results
      y_pred = rfrmodel.predict(X_test)

      # calculating the mean squared error
      MSE = mean_squared_error(y_test, y_pred)
      print('MSE: ', MSE)

      # Calculating the root mean squared error
```

```

RMSE = sqrt(mean_squared_error(y_test, y_pred))
print('RMSE: ', RMSE)

# Calculating the r2 score
R2 = r2_score(y_test, y_pred)
print('R2 :', R2)

```

```

MSE:  0.1714123990163851
RMSE:  0.41401980510162206
R2 : 0.8386405991096147

```

```

[100]: from sklearn.neighbors import KNeighborsRegressor
        from sklearn.metrics import mean_squared_error, accuracy_score
        from math import sqrt
        import matplotlib.pyplot as plt
        %matplotlib inline

```

```

[102]: rmse_val = []                                #To store rmse values for different k
        for K in range(50):
            K = K + 1
            model = KNeighborsRegressor(n_neighbors=K)

            model.fit(X_train, y_train)              #fit the model
            y_pred = model.predict(X_test)           #prediction on Test set
            error = sqrt(mean_squared_error(y_test, y_pred))    #calculate rmse
            rmse_val.append(error)                   #store rmse values
            print('RMSE value for K= ', K, ':is', error)

```

```

RMSE value for K=  1 :is 1.1200245101685928
RMSE value for K=  2 :is 0.9834698043100614
RMSE value for K=  3 :is 0.8690500411222918
RMSE value for K=  4 :is 0.8257511222522611
RMSE value for K=  5 :is 0.7773336566937904
RMSE value for K=  6 :is 0.7515511968123927
RMSE value for K=  7 :is 0.746304822381118
RMSE value for K=  8 :is 0.743420775790797
RMSE value for K=  9 :is 0.7485157220181602
RMSE value for K= 10 :is 0.7490477843371958
RMSE value for K= 11 :is 0.7578551544831839
RMSE value for K= 12 :is 0.7587120793465382
RMSE value for K= 13 :is 0.7687810594427169
RMSE value for K= 14 :is 0.7769897660229889
RMSE value for K= 15 :is 0.7840925893922719
RMSE value for K= 16 :is 0.7754834756596681
RMSE value for K= 17 :is 0.7848044400742131
RMSE value for K= 18 :is 0.7953809658266138
RMSE value for K= 19 :is 0.799694231540227
RMSE value for K= 20 :is 0.8012282411541236


```

```
RMSE value for K= 21 :is 0.8013957903530937
RMSE value for K= 22 :is 0.7954968968610465
RMSE value for K= 23 :is 0.7984572371131838
RMSE value for K= 24 :is 0.8012800992430992
RMSE value for K= 25 :is 0.8047112444891865
RMSE value for K= 26 :is 0.8076665712132023
RMSE value for K= 27 :is 0.8076223039795957
RMSE value for K= 28 :is 0.8102799556316295
RMSE value for K= 29 :is 0.8154113221651668
RMSE value for K= 30 :is 0.8161413870742454
RMSE value for K= 31 :is 0.8138198515624838
RMSE value for K= 32 :is 0.8186373042289483
RMSE value for K= 33 :is 0.818906336228432
RMSE value for K= 34 :is 0.817774613402975
RMSE value for K= 35 :is 0.8179928782407337
RMSE value for K= 36 :is 0.8211355439457957
RMSE value for K= 37 :is 0.8189292813072588
RMSE value for K= 38 :is 0.8232117326098805
RMSE value for K= 39 :is 0.8264046290972868
RMSE value for K= 40 :is 0.8246009409339825
RMSE value for K= 41 :is 0.8259723856115664
RMSE value for K= 42 :is 0.8254412335525577
RMSE value for K= 43 :is 0.8282410004927719
RMSE value for K= 44 :is 0.8267339842671623
RMSE value for K= 45 :is 0.8261882627121689
RMSE value for K= 46 :is 0.8246532346906921
RMSE value for K= 47 :is 0.8249664407992714
RMSE value for K= 48 :is 0.8268675384616488
RMSE value for K= 49 :is 0.8255957286306069
RMSE value for K= 50 :is 0.8217556915447253
```

```
[103]: # Plot the Graph for K Values

curve = pd.DataFrame(rmse_val) # Elbow Curve
curve.plot()
```

```
[103]: <AxesSubplot:>
```



```
[104]: from sklearn.model_selection import GridSearchCV

params = {'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]}

knn = KNeighborsRegressor()

model1 = GridSearchCV(knn, params, cv=5)
model1.fit(X_train, y_train)
model1.best_params_
```


```
[104]: {'n_neighbors': 8}
```

```
[105]: A2.columns
```

```
[105]: Index(['Battery_Capacity_kWh', 'Motor_Power_kW', 'Weight_kg',
          'Temperature_Celsius', 'Distance_Traveled_km', 'Average_Speed_kmph',
          'Energy_Consumption_kWh_per_100km', 'Highway', 'Suburban', 'Urban',
          'Fast', 'Home', 'Public'],
          dtype='object')
```


```
[106]: A2.hist(column='Battery_Capacity_kWh')
```

```
[106]: array([[<AxesSubplot:title={ 'center': 'Battery_Capacity_kWh' }>]],
          dtype=object)
```



```
[108]: #Scatter plots between pm2.5 and pressure.  
plt.scatter(x=A2['Motor_Power_kW'],y=A2['Distance_Traveled_km'])
```

```
[108]: <matplotlib.collections.PathCollection at 0x2319731b148>
```



[]: